



smoothing



- One drawback of word counting is that the number of parameters to estimate in the joint distribution over words and labels grows quadratically with the size of the vocabulary and label set. Thus many word-label pairs are never seen in the training set (Zipf's law makes this scaling even worse). One way to address this issue is via smoothing; which reduces the estimators variance at the expense of increasing its bias.

$$\hat{\phi}_{y,j} = \frac{\alpha + \sum_{i: y^{(i)}=y}^M x_j^{(i)}}{Na + \sum_{j'=1}^N \sum_{i: y^{(i)}=y}^M x_{j'}}$$

where  $\alpha$  is the smoothing parameter

# Smoothing

- One drawback of word counting is that the number of parameters to estimate in the joint distribution over words and labels grows quadratically with the size of the vocabulary and label set. Thus many word-label pairs are never seen in the training set (Zipf's law makes this scaling even worse). One way to address this issue is via smoothing; which reduces the estimators variance at the expense of increasing its bias.

$$\hat{\phi}_{y,j} = \frac{\alpha + \sum_{i: y^{(i)}=y}^M x_j^{(i)}}{N\alpha + \sum_{j'=1}^N \sum_{i: y^{(i)}=y}^M x_{j'}}$$

where  $\alpha$  is the smoothing parameter

# Naive Bayes' classifier example w/Laplace smoothing

	Cat	Documents
Training	-	just plain boring
	-	entirely predictable and lacks energy
	-	no surprises and very few laughs
	+	very powerful
	+	the most fun film of the summer
Test	?	predictable with no fun

**function** TRAIN NAIVE BAYES(D, C) **returns**  $\log P(c)$  and  $\log P(w|c)$

**for each** class  $c \in C$  # Calculate  $P(c)$  terms

$N_{doc}$  = number of documents in D

$N_c$  = number of documents from D in class c

$\logprior[c] \leftarrow \log \frac{N_c}{N_{doc}}$

$V \leftarrow$  vocabulary of D

$bigdoc[c] \leftarrow$  **append**(d) **for** d  $\in$  D **with** class c

**for each** word w in V # Calculate  $P(w|c)$  terms

$count(w, c) \leftarrow$  # of occurrences of w in  $bigdoc[c]$

$\loglikelihood[w, c] \leftarrow \log \frac{count(w, c) + 1}{\sum_{w' \in V} (count(w', c) + 1)}$

**return**  $\logprior$ ,  $\loglikelihood$ , V

**function** TEST NAIVE BAYES( $testdoc$ ,  $\logprior$ ,  $\loglikelihood$ , C, V) **returns** best c

**for each** class  $c \in C$

$sum[c] \leftarrow \logprior[c]$

**for each** position i in  $testdoc$

$word \leftarrow testdoc[i]$

**if** word  $\in V$

$sum[c] \leftarrow sum[c] + \loglikelihood[word, c]$

**return**  $\operatorname{argmax}_c sum[c]$